

Efficient model Inference

ChatGPT 3.5 ▾



How can I help you today?

Design a database schema
for an online merch store

Help me study
vocabulary for a college entrance exam

Tell me a fun fact
about the Roman Empire

Plan an itinerary
for a fashion-focused exploration of Paris

Message ChatGPT...



ChatGPT can make mistakes. Consider checking important information.



ĐẠI HỌC
HANOI UNIVERSITY

Time To First Token

Context decoding



Tokens Per Second



You

Tell me a random fun fact about the Roman Empire



ChatGPT

Sure! Did you know that the ancient Romans used urine as a cleaning agent? They believed that the ammonia in urine could help whiten and brighten their clothes. They even col ●

Generation step

Tokens Per Second

TPS for single query

TPS for single instance

Assistant

40 TPS is OK (or is it?)

Economics


Offline

Do we care?


Most efficient way

Inference



- Prompt processing
- Autoregressive steps
- Hardware utilization




- Serving queries
- Business cases




LLM response




Initiation phase




Decoding phase




Decoder block




Initiation phase




Generation step #1




Generation step #2




KV cache



Batch




Batch




Tesla V100, float16, maximum sequence length=64, average sequence length≈40

batch_size	XLA (in ms)	Faster Transformer (in ms)	Speedup over XLA	Effective Transformer (in ms)	Speedup over XLA
100	28.31	20.27	1.40	16.03	1.77
200	54.47	40.08	1.36	30.15	1.81
300	80.53	59.11	1.36	41.27	1.95
400	106.5	78.38	1.36	54.12	1.97
500	132.35	98.03	1.37	65.92	2.01
1000	261.18	190.91	1.38	133.61	1.95

Continuous batch




Flash attention



$$\mathbf{O} = \text{Dropout}(\text{Softmax}(\text{Mask}(\mathbf{QK}^T)))\mathbf{V}$$


Naive implementation requires
repeated R/W from slow GPU HBM.
Hard to scale to long sequences

Flash attention




$$O = \text{Dropout}(\text{Softmax}(\text{Mask}(\mathbf{QK}^T)))\mathbf{V}$$

Naive implementation requires
repeated R/W from slow GPU HBM.
Hard to scale to long sequences



Flash attention




Flash attention

Decomposing large softmax into smaller ones by scaling.



$$\text{softmax}([A_1, A_2]) = [\alpha \text{softmax}(A_1), \beta \text{softmax}(A_2)]$$

$$\text{softmax}([A_1, A_2]) \begin{bmatrix} V_1 \\ V_2 \end{bmatrix} = \alpha \text{softmax}(A_1) V_1 + \beta \text{softmax}(A_2) V_2$$

1. Load inputs by blocks from HBM to SRAM.
2. On chip, compute attention output wrt that block.
3. Update output in HBM by scaling.



Flash attention



2-4x speedup — with no approximation!

10-20x memory reduction — memory linear in sequence length

Flash attention 2






Figure 5: Attention forward speed on A100 GPU


Flash decoding




Flash decoding



Flash decoding



What about long sequence batches?



Paged Attention

0. Before generation.

Seq A Prompt: "Alan Turing is a computer scientist"
Completion: ""

Logical KV cache blocks

Block 0				
Block 1				
Block 2				
Block 3				

Block table


Physical block no.	# Filled slots
—	—
—	—
—	—
—	—

Physical KV cache blocks


Block 0				
Block 1				
Block 2				
Block 3				
Block 4				
Block 5				
Block 6				
Block 7				

Paged Attention

0. Shared prompt: Map logical blocks to the same physical blocks.



Tensor parallel



Performance of GPT-20B

Batch_size	Input Length	Output Length	Latency of single GPU (ms)	Latency of 2-way TP (ms)	Latency of 4-way TP (ms)	Latency of 8-way TP (ms)
1	20	8	225	147	102	89
2	20	8	225	152	108	94
4	20	8	228	158	113	100
8	20	8	239	169	121	107
16	20	8	268	191	133	113
32	20	8	331	230	155	127
64	20	8	452	314	200	169
128	20	8	726	484	318	256
256	20	8	1352	844	533	416
1	60	20	560	358	248	212
2	60	20	562	378	262	222
4	60	20	582	393	274	236
8	60	20	635	429	299	247
16	60	20	748	510	345	272
32	60	20	933	620	418	325
64	60	20	1352	887	574	454
128	60	20	2218	1384	928	699
256	60	20	4141	2424	1574	1152




Frameworks


- <https://github.com/vllm-project/vllm>
<https://github.com/InternLM/lndeploy>
<https://github.com/microsoft/DeepSpeed-MII>
<https://github.com/NVIDIA/TensorRT-LLM>
- <https://github.com/ggerganov/llama.cpp>

Pareto Curve: Quality vs Compute


Latency/RPS on the X-axis
Quality on the Y-axis



Compute bound vs Memory bound




Throughput = Y




Throughput?
 $1024 * Y$?

Compute bound vs Memory bound



$$\text{FLOP} = 10^9$$
$$\text{Bandwidth} = 2 * 10^6$$




$$\text{FLOP} = 10^6$$
$$\text{Bandwidth} = 10^6$$

FLOPS vs Bandwidth utilisation

MBU = (achieved memory bandwidth) / (peak memory bandwidth)

achieved memory bandwidth = (total model parameter size + KV cache size) / TPOT



For example:

7B model in float16 has TPOT = 14ms

14 GB parameters in 14ms => 1TB/s bandwidth

peak bandwidth = 2TB / sec

MBU=50%

Knowledge Distillation General framework

- **Teacher $p(y|x)$** : usually a LLM, e.g. GPT-3 (175B) achieves SOTA quality
- **Student $q(y|x)$** : small LM, e.g. T5 XL (3B) unable to reach teacher's quality by ordinary training
- **Knowledge Distillation (KD)**: process of teaching the student to imitate teacher's performance

$$L(p(y|x), q_{\theta}(y|x)) \rightarrow \min_{\theta}$$

Hard-label Distillation

- Sample pairs from teacher
- Finetune model on that pairs

$$\mathbb{E}_{p(y)} \log q_{\theta}(y) \rightarrow \max_{\theta}$$

$$y^{(1)}, \dots, y^{(N)} \sim p(y)$$

$$\frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \log q_{\theta}(y_t^{(n)} | y_{<t}^{(n)})$$

Soft-label Distillation

- Same as hard-label, but reproduce logits
- Sample from teacher by default
- Hack: use targets from original dataset and compute logits in parallel

$$\frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \sum_{v \in \mathcal{V}} p(y_t^{(n)} = v | y_{<t}^{(n)}) \log q_{\theta}(y_t^{(n)} = v | y_{<t}^{(n)})$$

KL Distillation: Idea


- Distance between distributions
- Monte-Carlo estimation

$$D_{\text{KL}}(p(y) \parallel q_{\theta}(y)) \rightarrow \min_{\theta}$$

$$\mathbb{E}_{p(y)} \log \frac{p(y)}{q_{\theta}(y)} = -\mathbb{E}_{p(y)} \log q_{\theta}(y) + \text{const} \rightarrow \min_{\theta}$$

Tiny Bert


- $L_{attn} = \frac{1}{h} \sum MSE(A_i^S, A_i^T)$
- $L_{hidden} = MSE(H^S W_h, H^T)$
- $L_{embed} = MSE(E^S W_e, E^T)$
- $L_{pred} = CE(\frac{z^T}{t}, \frac{z^S}{t})$
- H^S, H^T - хиддены студента и учителя
- A^S, A^T - аттеншн мапы



KL Distillation: LLM issue

- Student's distribution should cover teacher's distribution
- While naturally student is less expressive than teacher

$$D_{\text{KL}}(p(y) \parallel q_{\theta}(y)) = \mathbb{E}_{p(y)} \log \frac{p(y)}{q_{\theta}(y)}$$



- Swap arguments of KL
- Student approx. only the top probs

$$D_{\text{KL}}(q_{\theta}(y) \parallel p(y)) = \mathbb{E}_{q_{\theta}(y)} \log \frac{q_{\theta}(y)}{p(y)}$$

- Swap arguments of KL
- Student approx. only the top probs
- **Another problem occurs!**
- Unable to differentiate by sampled y

$$D_{\text{KL}}(q_{\theta}(y) \parallel p(y)) = \mathbb{E}_{q_{\theta}(y)} \log \frac{q_{\theta}(y)}{p(y)}$$

$$\tilde{q}_{\theta}(y_t \mid y_{<t}) = \alpha p(y_t \mid y_{<t}) + (1 - \alpha) q_{\theta}(y_t \mid y_{<t})$$

$$R_{t+1}^{\text{Norm}} = \frac{1}{T - t - 1} \sum_{k=t}^T \log \frac{q_{\theta}(y_k \mid y_{<k})}{p(y_k \mid y_{<k})}$$

- Importance sampling
- Regularisation term
- Length normalisation

$$\begin{aligned} \nabla_{\theta} \mathcal{L} = \mathbb{E}_{\tilde{q}_{\theta}(y)} \sum_{t=1}^T w_t & \left[R_{t+1}^{\text{Norm}} \nabla_{\theta} \log q_{\theta}(y_t \mid y_{<t}) \right. \\ & \left. + \nabla_{\theta} \mathbb{E}_{q_{\theta}(y_t \mid x, y_{<t})} \log \frac{q_{\theta}(y_t \mid y_{<t})}{p(y_t \mid y_{<t})} \right] \end{aligned}$$

SLIM: Alternative

- Take only top %5 logits
- Compatibility with CERL (behaviour cloning)


$$L_{final} = L_{ce} + \alpha \left(1 - \exp\left(-\frac{s_i}{t_i}\right)\right) L_{kd}$$

LLaMA						
	DollyEval		VicunaEval		SelfInst	
	Rouge-L	GPT-4	Rouge-L	GPT-4	Rouge-L	GPT-4
Teacher (LLaMA 13B)	29.7	85.3	19.4	66.3	23.4	76.2
SFT	26.3	79.47	17.5	61.5	20.8	70.6
MiniLLM	28.7	82	19.8	63.8	20.5	72.1
SLIM (Ours)	29.2	82.6	20.1	64.1	23.2	73.2

LLaMA 2						
	DollyEval		VicunaEval		SelfInst	
	Rouge-L	GPT-4	Rouge-L	GPT-4	Rouge-L	GPT-4
Teacher (LLaMA 2 13B)	30.2	88.9	21.3	69.3	25.1	79.1
SFT	26.5	80.3	18.3	62.8	21.3	73.4
SLIM (Ours)	29.3	84.6	19.9	67.1	23.4	75.1

MPT						
	DollyEval		VicunaEval		SelfInst	
	Rouge-L	GPT-4	Rouge-L	GPT-4	Rouge-L	GPT-4
Teacher (MPT-30B-instruct)	44.0	94.7	19.3	67.3	23.5	76.1
SFT	28.6	79.5	16.62	60.7	19.9	70.7
SLIM (Ours)	31.1	83.3	17.83	63.4	22.9	73.8

Table 1: We report the Rouge-L and GPT-4 agreement scores on 3 different datasets across 3 different models. We do not have MiniLLM numbers for LLaMA 2 and MPT experiments since the authors did not open-source their models with these backbones.



- Hard label KD is all you need
- RL as distillation very good

- FP16 -> INT 1/2/4/8
- VRAM and latency/rps boost
- WxAy
- Symmetric vs Asymmetric
- Granularity
- Grid


Quantizing a real-valued tensor $\mathbf{x}$ is performed by first mapping it to an unsigned integer grid:

$$\mathbf{x}^{(\mathbb{Z})} = \text{clip} \left(\left\lfloor \frac{\mathbf{x}}{s} \right\rfloor + z; 0, 2^b - 1 \right), \quad (1)$$

It is possible to approximately recover the real-valued input $\mathbf{x}$ through an operation that is often referred to as *de-quantization*:

$$\hat{\mathbf{x}} := q(\mathbf{x}; s, z, b) = s \left(\mathbf{x}^{(\mathbb{Z})} - z \right) \approx \mathbf{x}. \quad (2)$$

In the case of symmetric quantization, we restrict the quantization grid to be symmetric around z .




PTQ

- No/~500/1000 calibration samples
- <10 GPU hours
- Calibrating s(scale), z(zero-point)

QAT

- Huge dataset
- Better quality
- Hard to implement



Quality drop

- Bert example
- Weights - easy
- Activations - hard
- Outliers - problem

Configuration	CoLA	SST-2	MRPC	STS-B	QQP	MNLI	QNLI	RTE	GLUE
FP32	57.27	93.12	88.36	89.09	89.72	84.91	91.58	70.40	83.06
W8A8	54.74	92.55	88.53	81.02	83.81	50.31	52.32	64.98	71.03
W32A8	56.70	92.43	86.98	82.87	84.70	52.80	52.44	53.07	70.25
W8A32	58.63	92.55	88.74	89.05	89.72	84.58	91.43	71.12	83.23

Table 1: Post-training quantization results on development sets of the GLUE benchmark (except WNLI). The metrics for these tasks can be found in the GLUE paper (Wang et al., 2018a); in all cases, higher is better. FP32 baseline is trained by the authors from the pre-trained checkpoint, see Appendix B.1 for details. We report a median over 5 runs with different random seeds.

Quantized activations	STS-B	MNLI	QNLI	RTE
none (FP32 model)	89.09	84.91	91.58	70.40
all	62.64	42.67	50.74	48.74
all, except softmax input	70.92	42.54	51.84	48.74
all, except sum of embeddings	67.57	46.82	51.22	51.26
all, except self-attention output	70.47	46.57	50.98	50.90
all, except softmax output	72.83	50.35	50.23	49.46
all, except residual connections after FFN	81.57	82.56	89.73	67.15
same as above, but for layers 10, 11 only	79.40	81.24	88.03	63.90








Table 2 & Figure 1: *Left*: Leave-one-out analysis for activation quantizers on problematic GLUE tasks. We set all weights to FP32 and use current min-max (with a batch size of 1) range estimator for activations. We report median score over 5 runs with different random seeds. *Right*: A schematic illustration of the attention layer in BERT. Hidden activation tensor is denoted by x . $\oplus$ is an element-wise addition. A problematic residual connection sum after feed-forward network is highlighted in red.


Outliers is problem

- Outliers occurs in QKVO projections
- Outliers have huge impact on quality



Intuition


- Migrate some difficulty to weights



- W8A8 with same quality
- Apply to all BMM
- Fuse scaling with previous operations

Table 2: Among different activation quantization schemes, only per-channel quantization (Bondarenko et al., 2021) preserves the accuracy, but it is *not* compatible (marked in gray) with INT8 GEMM kernels. We report the average accuracy on WinoGrande, HellaSwag, PIQA, and LAMBADA.

Model size (OPT-)	6.7B	13B	30B	66B	175B
FP16	64.9%	65.6%	67.9%	69.5%	71.6%
INT8 per-tensor	39.9%	33.0%	32.8%	33.1%	32.3%
INT8 per-token	42.5%	33.0%	33.1%	32.9%	31.7%
INT8 per-channel	64.8%	65.6%	68.0%	69.4%	71.4%



Method

$$X^{INT8} = \left[\frac{X^{FP16}}{s} \right], s = \frac{\max(|X|)}{127}$$

$$Y = (X \text{diag}(s)^{-1}) \cdot (\text{diag}(s)W)$$

$$s_j = \frac{\max(|X_j|)^a}{\max(|W_j|)^{1-a}}$$




Table 8: GPU Latency (ms) of different quantization schemes. The coarser the quantization scheme (from per-token to per-tensor, dynamic to static, O1 to O3, defined in Table 3), the lower the latency. SmoothQuant achieves lower latency compared to FP16 under all settings, while LLM.int8() is mostly slower. The batch size is 4.

Model	OPT-13B		OPT-30B	
	Sequence Length		Sequence Length	
	256	512	256	512
FP16	152.6	296.3	343.0	659.9
LLM.int8()	237.1	371.5	387.9	654.9
SmoothQuant-O1	124.5	243.3	246.7	490.7
SmoothQuant-O2	120.5	235.1	240.2	478.3
SmoothQuant-O3	112.1	223.1	227.6	458.4





Figure 7: The PyTorch implementation of SmoothQuant-O3 achieves up to $1.51\times$ speedup and $1.96\times$ memory saving for OPT models on a single NVIDIA A100-80GB GPU, while LLM.int8() slows down the inference in most cases.

GPU	FP16	3bit	Speedup	GPU reduction
A6000 – 48GB	589ms	130ms	4.53×	8 → 2
A100 – 80GB	230ms	71ms	3.24×	5 → 1

Table 6: Average per-token latency (batch size 1) when generating sequences of length 128.

- Decoding only
- Published code and CUDA kernels
- Solving minimization task layer per layer

$$\operatorname{argmin}_{\hat{W}} ||WX - \hat{W}X||_2^2$$



Method	Bits	OPT-175B				BLOOM-176B			
		Wiki2	PTB	C4	LAMB. ↑	Wiki2	PTB	C4	LAMB. ↑
Baseline	16	8.34	12.01	10.13	75.59	8.11	14.59	11.71	67.40
RTN	4	10.54	14.22	11.61	71.34	8.37	15.00	12.04	66.70
GPTQ	4	8.37	12.26	10.28	76.80	8.21	14.75	11.81	67.71
RTN	3	7.3e3	8.0e3	4.6e3	0	571.	107.	598.	0.17
GPTQ	3	8.68	12.68	10.67	76.19	8.64	15.57	12.27	65.10
GPTQ	3/g1024	8.45	12.48	10.47	77.39	8.35	15.01	11.98	67.47
GPTQ	3/g128	8.45	12.37	10.36	76.42	8.26	14.89	11.85	67.86

Table 5: Results summary for OPT-175B and BLOOM-176B. “g1024” and “g128” denote results with groupings of size 1024 and 128, respectively.


Architecture

Encoder-decoder vs decoder


Who faster with same amount of parameters ?




Bert tips




(a) Cross-Attention Model (e.g., BERT reranker)



(b) Dense Retrievers (e.g., DPR)






(c) ColBERT: All-to-All Match



(d) COIL: Contextualized Exact Match

Bert tips



Speculative decoding: Vanilla idea

- Small draft model, e.g 7B
- Large verification model, e.g. 70B
- Decode 1 to $K+1$ tokens per iteration
- Verify by large model

Table 1 | Chinchilla performance and speed on XSum and HumanEval with naive and speculative sampling at batch size 1 and $K = 4$. XSum was executed with nucleus parameter $p = 0.8$, and HumanEval with $p = 0.95$ and temperature 0.8.

Sampling Method	Benchmark	Result	Mean Token Time	Speed Up
ArS (Nucleus)	XSum (ROUGE-2)	0.112	14.1ms/Token	1×
SpS (Nucleus)		0.114	7.52ms/Token	1.92×
ArS (Greedy)	XSum (ROUGE-2)	0.157	14.1ms/Token	1×
SpS (Greedy)		0.156	7.00ms/Token	2.01×
ArS (Nucleus)	HumanEval (100 Shot)	45.1%	14.1ms/Token	1×
SpS (Nucleus)		47.0%	5.73ms/Token	2.46×

Algorithm 2 Speculative Sampling (SpS) with Auto-Regressive Target and Draft Models

Given lookahead K and minimum target sequence length T .

Given auto-regressive target model $q(.|.)$, and auto-regressive draft model $p(.|.)$, initial prompt sequence $x_0, \dots, x_t$.

Initialise $n \leftarrow t$.

while $n < T$ **do**

for $t = 1 : K$ **do**

 Sample draft auto-regressively $\tilde{x}_t \sim p(x|x_1, \dots, x_n, \tilde{x}_1, \dots, \tilde{x}_{t-1})$

end for

 In parallel, compute $K + 1$ sets of logits from drafts $\tilde{x}_1, \dots, \tilde{x}_K$:

$$q(x|x_1, \dots, x_n), q(x|x_1, \dots, x_n, \tilde{x}_1), \dots, q(x|x_1, \dots, x_n, \tilde{x}_1, \dots, \tilde{x}_K)$$

for $t = 1 : K$ **do**

 Sample $r \sim U[0, 1]$ from a uniform distribution.

if $r < \min\left(1, \frac{q(x|x_1, \dots, x_{n+t-1})}{p(x|x_1, \dots, x_{n+t-1})}\right)$ **then**

 Set $x_{n+t} \leftarrow \tilde{x}_t$ and $n \leftarrow n + 1$.

else

 sample $x_{n+t} \sim (q(x|x_1, \dots, x_{n+t-1}) - p(x|x_1, \dots, x_{n+t-1}))_+$ and exit for loop.

end if

end for

 If all tokens $x_{n+1}, \dots, x_{n+K}$ are accepted, sample extra token $x_{n+K+1} \sim q(x|x_1, \dots, x_n, x_{n+K})$ and set $n \leftarrow n + 1$.

end while

Inference with reference

Algorithm 1 LLMA Decoding.


Input: $x, D = (d_1, \dots, d_n), n, k, N$;

Output: y ;

```

1:  $y \leftarrow []$ 
2: while  $\text{LEN}(y) < N$  do
3:    $\text{matched}, d, \text{pos} \leftarrow \text{MATCH\_NGRAMS}(y, D, n)$ 
4:   if  $\neg \text{matched}$  then
5:      $o = \text{LLM}(x, y)$ 
6:      $\text{APPEND}(y, o)$ 
7:     continue
8:   end if
9:    $(o_0, o_1, \dots, o_k) \leftarrow \text{LLM}(x, y, d_{\text{pos}}, \dots, d_{\text{pos}+k-1})$ 
10:   $\text{APPEND}(y, o_0)$ 
11:  for  $i$  in  $0, \dots, k-1$  do
12:    if  $o_i \neq d_{\text{pos}+i}$  then
13:      break
14:    end if
15:     $\text{APPEND}(y, o_{i+1})$ 
16:  end for
17: end while

```



Why speculative sampling is bad

- Optimize only MBU
- Work only with batch_size=1
- Don't optimize memory(otherwise)
- hard to implement

Conclusion

- Don't forget about MBU vs MFU!!!
- Smoothquant is universal, gptq/aqlm if you have small rps/need to host huge models on 1 GPU
- If you have reword model use it!, otherwise hard-label KD if you have a lot of compute, slim if not
- Don't forget about architecture optimizing